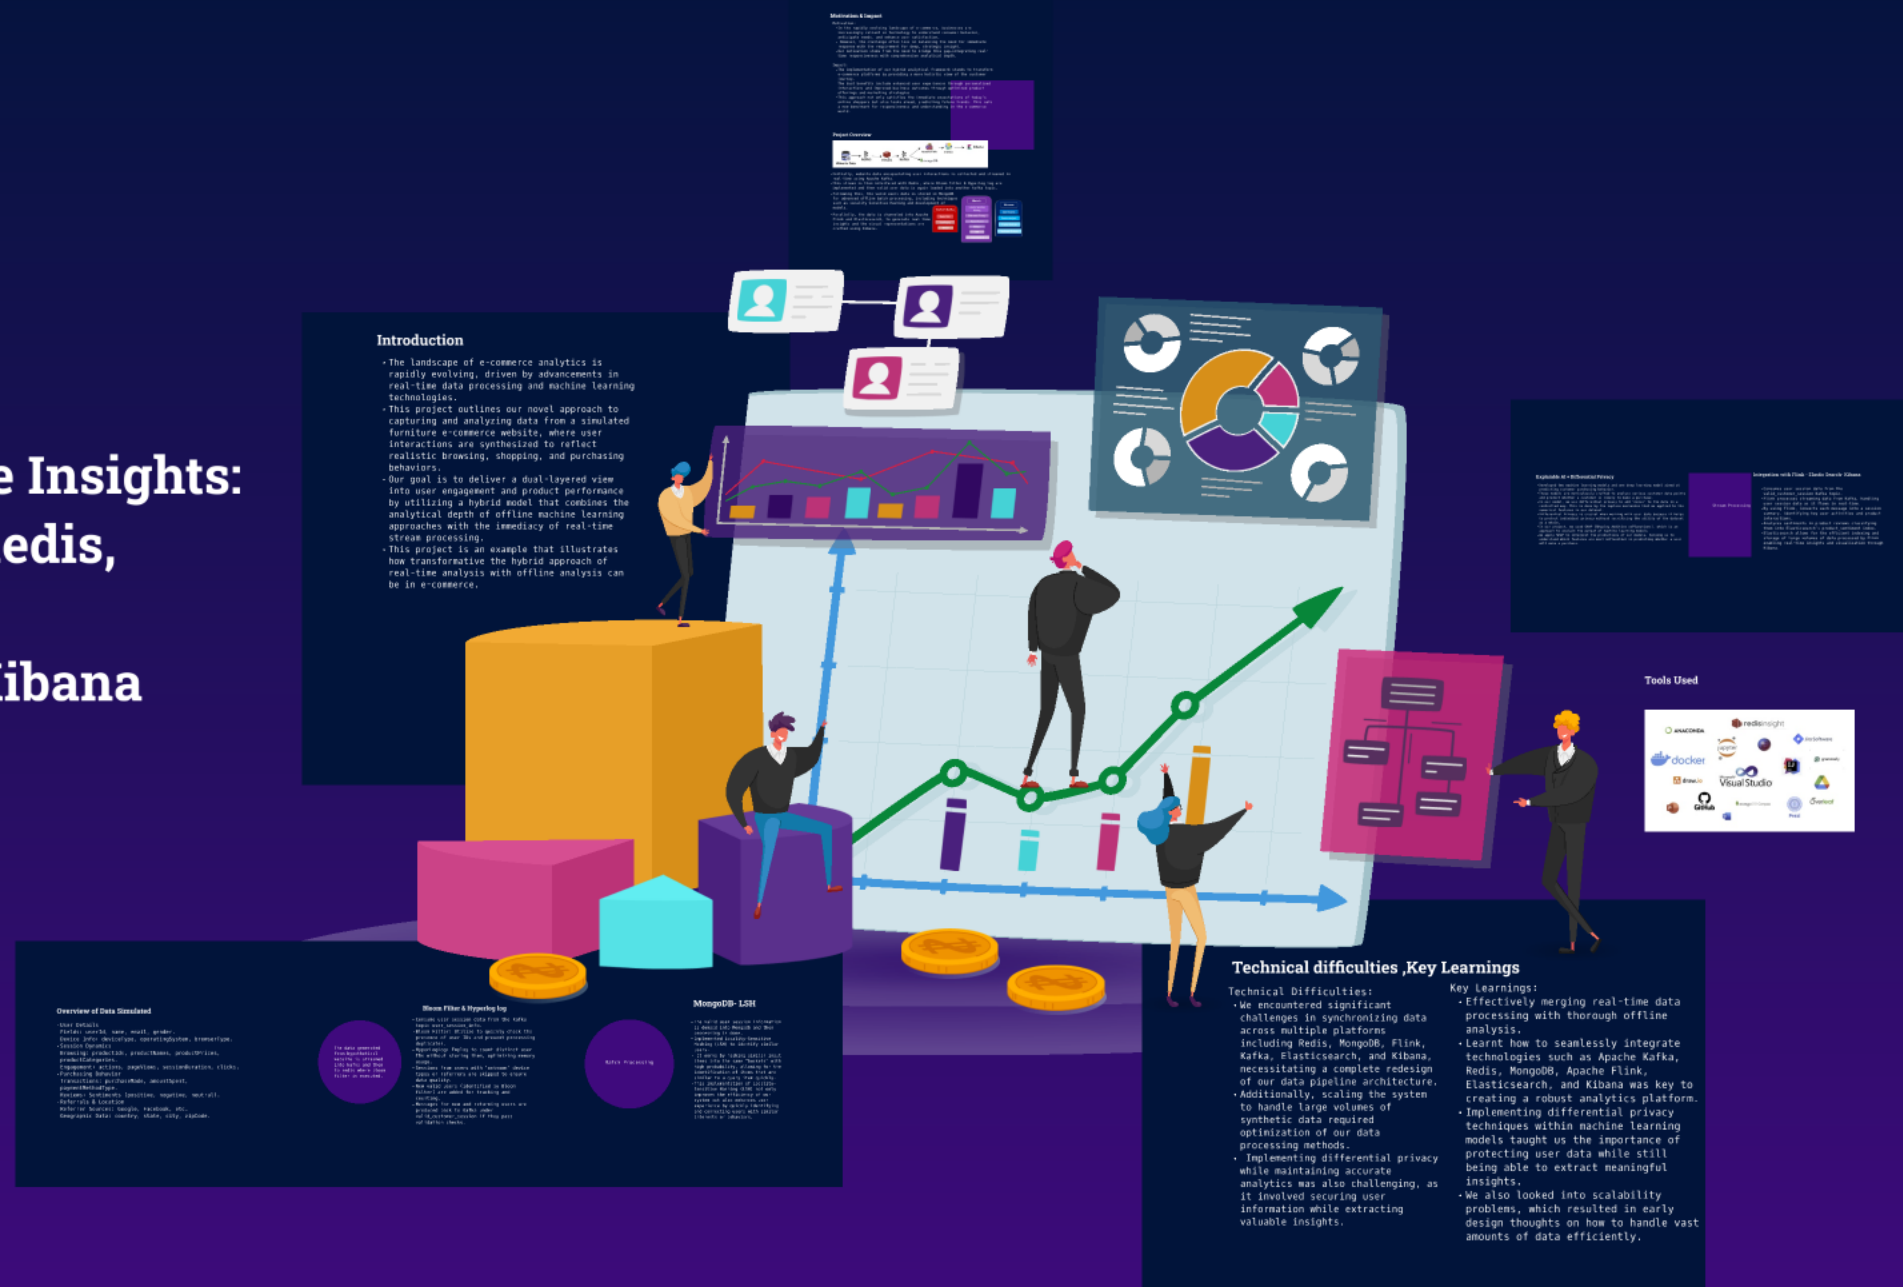


# Streamlined Website Insights: Leveraging Kafka, Redis, MongoDB, Flink, Elasticsearch, and Kibana

Group 2:

- Vidushi
- Kamakshi
- Akanksha
- Spandana



# Introduction

- The landscape of e-commerce analytics is rapidly evolving, driven by advancements in real-time data processing and machine learning technologies.
- This project outlines our novel approach to capturing and analyzing data from a simulated furniture e-commerce website, where user interactions are synthesized to reflect realistic browsing, shopping, and purchasing behaviors.
- Our goal is to deliver a dual-layered view into user engagement and product performance by utilizing a hybrid model that combines the analytical depth of offline machine learning approaches with the immediacy of real-time stream processing.
- This project is an example that illustrates how transformative the hybrid approach of real-time analysis with offline analysis can be in e-commerce.



# Motivation & Impact

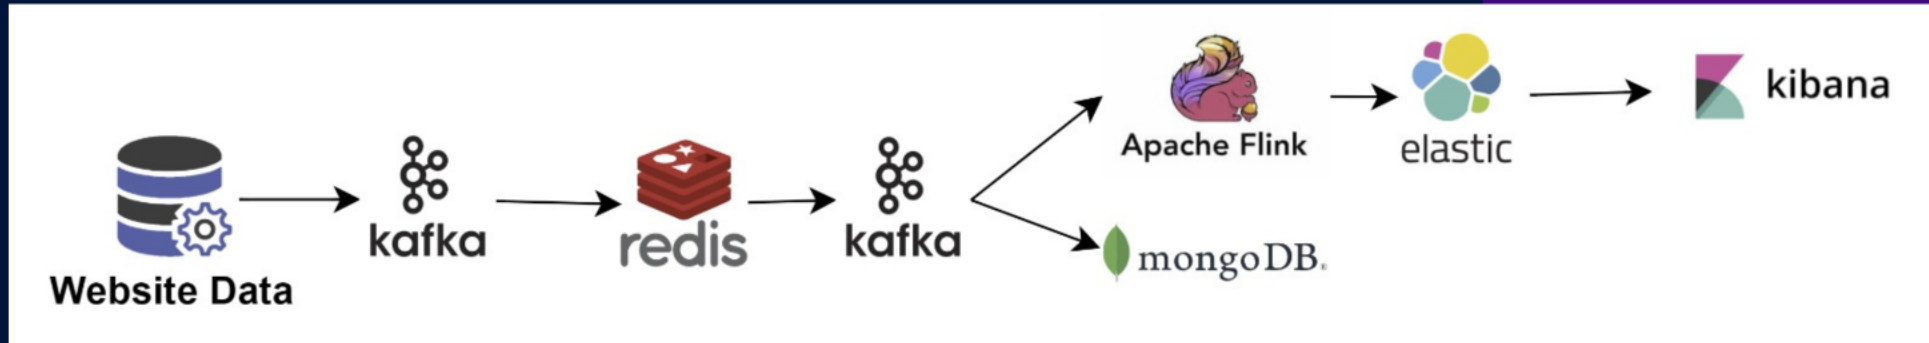
## Motivation:

- In the rapidly evolving landscape of e-commerce, businesses are increasingly reliant on technology to understand consumer behavior, anticipate needs, and enhance user satisfaction.
- However, the challenge often lies in balancing the need for immediate response with the requirement for deep, strategic insight.
- Our motivation stems from the need to bridge this gap—integrating real-time responsiveness with comprehensive analytical depth.

## Impact:

- The implementation of our hybrid analytical framework stands to transform e-commerce platforms by providing a more holistic view of the customer journey.
- The dual benefits include enhanced user experiences through personalized interactions and improved business outcomes through optimized product offerings and marketing strategies.
- This approach not only satisfies the immediate expectations of today's online shoppers but also looks ahead, predicting future trends. This sets a new benchmark for responsiveness and understanding in the e-commerce world.

# Project Overview



- Initially, website data encapsulating user interactions is collected and streamed in real-time using Apache Kafka.
- This stream is then interfaced with Redis, where Bloom Filter & Hyperlog log are implemented and then valid user data is again loaded into another kafka topic.
- Following this, the valid users data is stored in MongoDB for advanced offline batch processing, including techniques such as Locality Sensitive Hashing and development of models.
- Parallely, the data is channeled into Apache Flink and Elasticsearch, to generate real time insights and the visual representations are crafted using Kibana.

## Cache Checks

Bloom Filter

Hyperloglog

HINCRBY

## Batch

Locality Sensitive Hashing

Differential Privacy

RandomForest

XGBoost

ANN

Explainable AI

## Stream

User Analysis

Session Analysis

Product Analysis

Sentiment Analysis



# Overview of Data Simulated

- User Details

Fields: `userId`, `name`, `email`, `gender`.

Device Info: `deviceType`, `operatingSystem`, `browserType`.

- Session Dynamics

Browsing: `productIds`, `productNames`, `productPrices`,  
`productCategories`.

Engagement: `actions`, `pageViews`, `sessionDuration`, `clicks`.

- Purchasing Behavior

Transactions: `purchaseMade`, `amountSpent`,  
`paymentMethodType`.

Reviews: Sentiments (`positive`, `negative`, `neutral`).

- Referrals & Location

Referrer Sources: `Google`, `Facebook`, etc.

Geographic Data: `country`, `state`, `city`, `zipCode`.



## Bloom Filter & Hyperlog log

The data generated from hypothetical website is streamed into kafka and then to redis where bloom filter is executed.

- Consume user session data from the Kafka topic `user_session_info`.
- Bloom Filter: Utilize to quickly check the presence of user IDs and prevent processing duplicates.
- HyperLogLog: Employ to count distinct user IDs without storing them, optimizing memory usage.
- Sessions from users with 'unknown' device types or referrers are skipped to ensure data quality.
- New valid users (identified by Bloom filter) are added for tracking and counting.
- Messages for new and returning users are produced back to Kafka under `valid_customer_session` if they pass validation checks.

# MongoDB- LSH

Batch Processing

- The valid user session information is dumped into MongoDB and then processing is done.
- Implemented Locality-Sensitive Hashing (LSH) to identify similar users.
- It works by hashing similar input items into the same "buckets" with high probability, allowing for the identification of items that are similar to a query item quickly.
- This implementation of Locality-Sensitive Hashing (LSH) not only improves the efficiency of our system but also enhances user experience by quickly identifying and connecting users with similar interests or behaviors.





## Explainable AI + Differential Privacy

- Developed two machine learning models and one deep learning model aimed at predicting customer purchasing behavior.
- These models are meticulously crafted to analyze various customer data points and predict whether a customer is likely to make a purchase.
- In our model, we use differential privacy to add 'noise' to the data in a controlled way. This is done by the Laplace mechanism that we applied to the numerical features in our dataset.
- Differential Privacy is crucial when working with user data because it helps to protect individual privacy without sacrificing the utility of the dataset as a whole.
- In our project, we used SHAP (SHapley Additive exPlanations), which is an approach to explain the output of machine learning models.
- We apply SHAP to interpret the predictions of our models, helping us to understand which features are most influential in predicting whether a user will make a purchase.

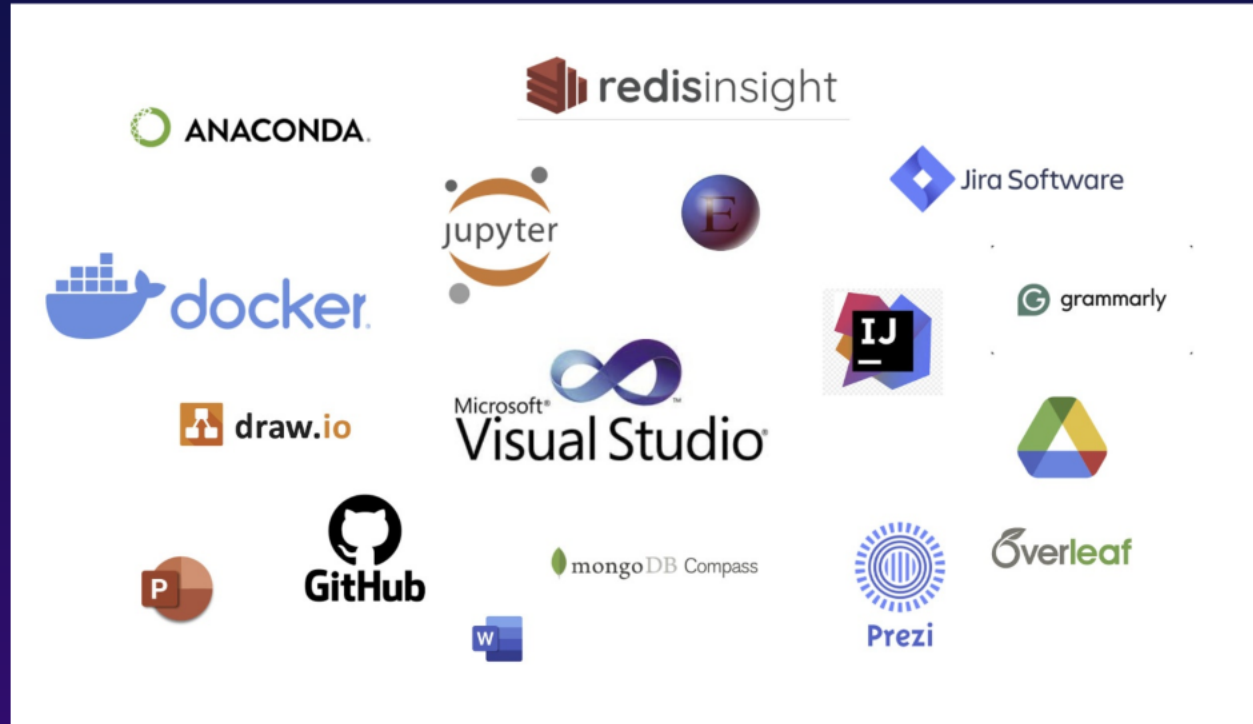
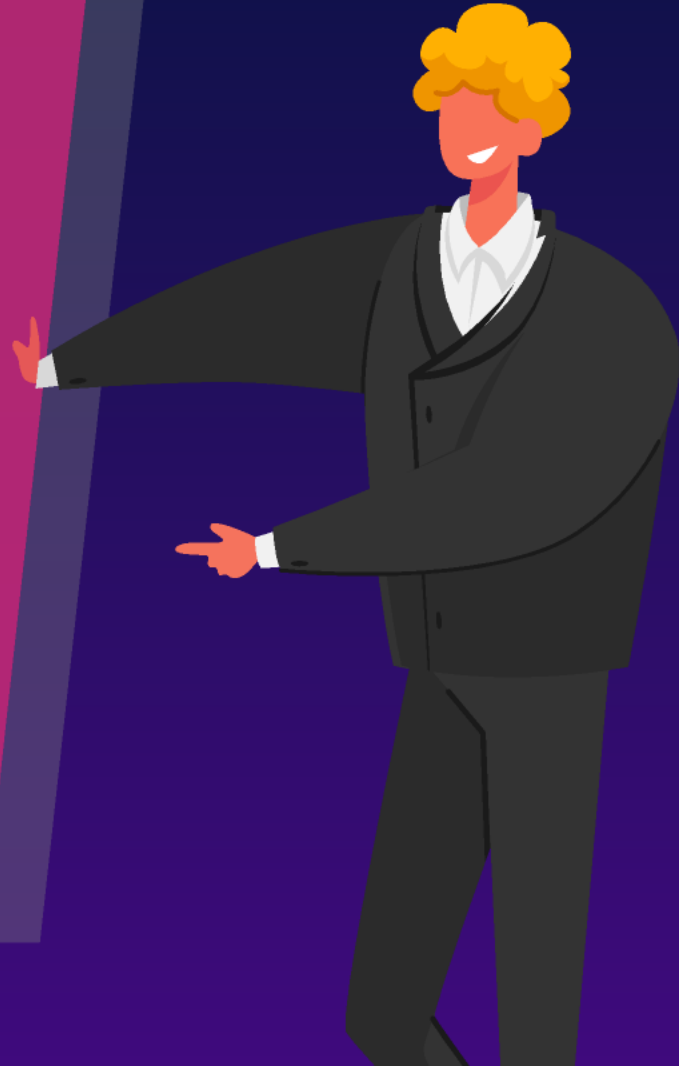
## Integration with Flink - Elastic Search- Kibana

- Consumes user session data from the `valid_customer_session` Kafka topic.
- Flink processes streaming data from Kafka, handling user session data as it flows in real-time.
- By using Flink, Converts each message into a session summary, identifying key user activities and product interactions.
- Analyzes sentiments in product reviews classifying them into Elasticsearch's `product_sentiment` index.
- Elasticsearch allows for the efficient indexing and storage of large volumes of data processed by Flink enabling real-time insights and visualization through Kibana.

Stream Processing



# Tools Used



# Technical difficulties ,Key Learnings

## Technical Difficulties:

- We encountered significant challenges in synchronizing data across multiple platforms including Redis, MongoDB, Flink, Kafka, Elasticsearch, and Kibana, necessitating a complete redesign of our data pipeline architecture.
- Additionally, scaling the system to handle large volumes of synthetic data required optimization of our data processing methods.
- Implementing differential privacy while maintaining accurate analytics was also challenging, as it involved securing user information while extracting valuable insights.

## Key Learnings:

- Effectively merging real-time data processing with thorough offline analysis.
- Learnt how to seamlessly integrate technologies such as Apache Kafka, Redis, MongoDB, Apache Flink, Elasticsearch, and Kibana was key to creating a robust analytics platform.
- Implementing differential privacy techniques within machine learning models taught us the importance of protecting user data while still being able to extract meaningful insights.
- We also looked into scalability problems, which resulted in early design thoughts on how to handle vast amounts of data efficiently.



